

# Query Optimization: A Library Adventure

*Using Storytelling, Hands-On Activities, Math & Diagrams to Learn SQL Tuning*

**Marc Reyes**

# Introduction – Imagine a Massive Library

- *Imagine you're searching for a book in a **huge library** with millions of books...*
- You need **"The Adventures of Data"** 📖. How will you find it?
  - **Option 1:** Wander aisle by aisle (😓 slow!)
  - **Option 2:** Use the library's **catalog** or ask a librarian (🚀 fast!)
- 🖱️ **Question:** *What would you do to find the book quickly?*

```
%%{init: {"theme": "default"}}%%  
flowchart LR  
  A["Start: Search for The Adventures of Data"]  
  B["Wander through every aisle (Full Table Scan)"]  
  C["Use Library Catalog (Index Lookup)"]  
  A --> B  
  A --> C
```

# The Library Analogy – Catalogs vs. Scanning Shelves

- **Library Catalog = Index:** Quickly finds books by title or author.
- **Scanning Every Shelf = Full Table Scan:** Checks every book one by one.
- A well-indexed database is like a **well-organized library!**
- *Using the catalog gets you there fast; scanning shelves is a last resort.*

```
%%{init: {"theme": "default"}}%%  
flowchart LR  
    A["Library Search"]  
    B["Using Catalog (Index)"]  
    C["Scanning Shelves (Full Scan)"]  
    A --> B  
    A --> C  
    B --> D["Fast Lookup"]  
    C --> E["Slow Search"]
```

# What is Query Optimization?

- **Query Optimization:** Finding the fastest path to your data.
- The **Query Optimizer** = the database's decision maker for search strategy.
- It evaluates options (scan vs. index, join methods, etc.) and picks the lowest-cost plan.
- *Goal:* Retrieve results **quickly and efficiently** 🚀

```
%%{init: {"theme": "default"}}%%  
flowchart TD  
    A["SQL Query"]  
    B["Parsing & Binding"]  
    C["Query Optimizer"]  
    D["Execution Engine"]  
    A --> B  
    B --> C  
    C --> D
```

# Database Query Optimization Landscape

- **Different Optimizers, Similar Principles:**
  - PostgreSQL: Cost-based optimizer with detailed statistics
  - MySQL: Hybrid rule/cost-based optimizer
  - Oracle: Advanced cost-based optimizer with adaptive plans
  - SQL Server: Cardinality estimation and statistics-driven
- **What they all share:** Finding the most efficient path to data

## Mathematical Modeling of Query Cost (1/2)

- **Sequential Scan Cost:**

$$\text{Cost}_{\text{seq}} = N \times C_{\text{read}}$$

- $N$ : Total rows,  $C_{\text{read}}$ : Cost per row read.

- **Index Scan Cost:**

$$\text{Cost}_{\text{index}} = \log_2(N) \times C_{\text{read}} + R \times C_{\text{read}}$$

- $R$ : Rows returned.

## Mathematical Modeling of Query Cost (2/2)

- **Nested Loop Join Cost:**

$$\text{Cost}_{\text{NLJ}} = N_{\text{outer}} \times N_{\text{inner}} \times C_{\text{compare}}$$

- **Hash Join Cost:**

$$\text{Cost}_{\text{hash}} \approx N_{\text{build}} + N_{\text{probe}}$$

- *These formulas help the optimizer decide which plan is cheaper.*

## Case Study – A Slow SQL Query (The Problem)

- **Scenario:** A `books` table with millions of records.
- **Query:**

```
SELECT *  
FROM books  
WHERE title = 'The Hobbit';
```

- **Question:** What happens when this runs on a huge table **without an index** on `title` ?
- **Answer:** A **full table scan**—it checks every row, which is painfully slow.



## Interactive Activity – Optimizing the Query (Brainstorm)

- **Your Turn:** How can we **speed up** this query?
- Ideas to consider:
  - **Add an index** on the `title` column.
  - Avoid `SELECT *` if only certain columns are needed.
  - Other strategies?
- **Share your ideas!** (*Let's hear a few guesses.*)

# The Climax – Applying the Optimization (Index to the Rescue!)

- **Solution:** Create an index on the `title` column:

```
CREATE INDEX idx_books_title ON books(title);
```

- Re-run the query:

```
SELECT * FROM books WHERE title = 'The Hobbit';
```

- **Result:**
  - **Before:** Full table scan (slow).
  - **After:** Index seek (fast!).
- *Performance improves dramatically!* 🚀

## Visual: Query Performance Before & After Optimization

```
%%{init: {"theme": "default"}}%%  
flowchart LR  
    subgraph "Query Execution Times"  
        Before["Before: 3500ms"] --->|"53x faster"| After["After: 65ms"]  
    end  
    style Before fill:#ffaana,color:#000  
    style After fill:#aafnaa,color:#000
```

*A typical database with millions of rows might see a 50x+ improvement!*

## Live Demo – Understanding the EXPLAIN Plan (Interactive)

- **Before Index:**
  - Execution plan: "Seq Scan on books" (scanning all rows).
- **After Index:**
  - Execution plan: "Index Scan using idx\_books\_title" (direct lookup).
- **Interactive:**
  - Spot the difference: "Scan" vs. "Index"
- The optimizer chose a new plan once the index was available!

## Live Demo – EXPLAIN Plan Diagram

```
%%{init: {"theme": "default"}}%%  
flowchart LR  
    A["SQL Query"]  
    B["Without Index"]  
    C["With Index"]  
    A --> B  
    A --> C  
    B --> D["Sequential Scan on books"]  
    C --> E["Index Scan using idx_books_title"]
```

## Real-World Query Optimization Impact

Company	Before Optimization	After Optimization	Improvement
E-commerce site	8.2 seconds	120 ms	68x faster
Social media app	12.5 seconds	350 ms	35x faster
Financial system	45 seconds	480 ms	93x faster

*Well-optimized queries save server resources, reduce costs, and improve user experience*

# Common Query Optimization Anti-Patterns

- Using functions on indexed columns:

```
-- Bad: Can't use index
SELECT * FROM users WHERE UPPER(email) = 'USER@EXAMPLE.COM';

-- Good: Can use index
SELECT * FROM users WHERE email = 'user@example.com';
```

- Unnecessary subqueries and complex joins
- Missing statistics or outdated indexes
- Wildcard at beginning of LIKE pattern: `WHERE name LIKE '%Smith'`

# Interactive Quiz: Spot the Optimized Query

Which query would execute faster?

## Query A:

```
SELECT * FROM customers
WHERE last_name LIKE '%son%'
ORDER BY created_at;
```

## Query B:

```
SELECT customer_id, first_name, last_name
FROM customers
WHERE customer_id BETWEEN 1000 AND 2000
ORDER BY customer_id;
```

*(Hint: Think about indexes and selectivity)*



# Key Takeaways – Tips for Query Optimization

- **Index Smartly:**
  - Use indexes on columns frequently searched or joined.
- **Be Selective:**
  - Avoid `SELECT *` ; fetch only needed columns.
- **Analyze with EXPLAIN:**
  - Understand and improve your query plans.
- **Think Like the Optimizer:**
  - Use cost estimates (like our formulas) to guide improvements.
- *Balance is key:* Too many indexes can hurt write performance.

## Optimization Strategies Beyond Indexing

- **Partitioning large tables** by date, region, or category
- **Materialized views** for complex aggregations
- **Query parameterization** to leverage plan caching
- **Using appropriate join types** (nested loop, hash, merge)
- **Data denormalization** for read-heavy scenarios
- **Query rewriting** to simplify complex logic

# Conclusion – The End of the Story

- **Recap:**
  - We transformed a slow query (wandering the library) into a fast one (using a catalog).
  - We applied math and practical tools to optimize performance.
- **Moral:**
  - Treat your database like a library—stay organized and use available tools.
- **Engage:**
  - Any questions or experiences to share?
- **Thank You & Happy Querying!** 🙌📊

# Resources to Learn More

- **Books:**
  - "SQL Performance Explained" by Markus Winand
  - "High Performance MySQL" by Schwartz, Zaitsev & Tkachenko
- **Online Tools:**
  - [EXPLAIN.DEPESZ.COM](https://explain.depesz.com) - PostgreSQL EXPLAIN visualizer
  - [Use The Index, Luke!](#) - SQL indexing tutorials
- **Documentation:**
  - Your DB's optimization guide (PostgreSQL, MySQL, etc.)
  - Query Plan analysis documentation