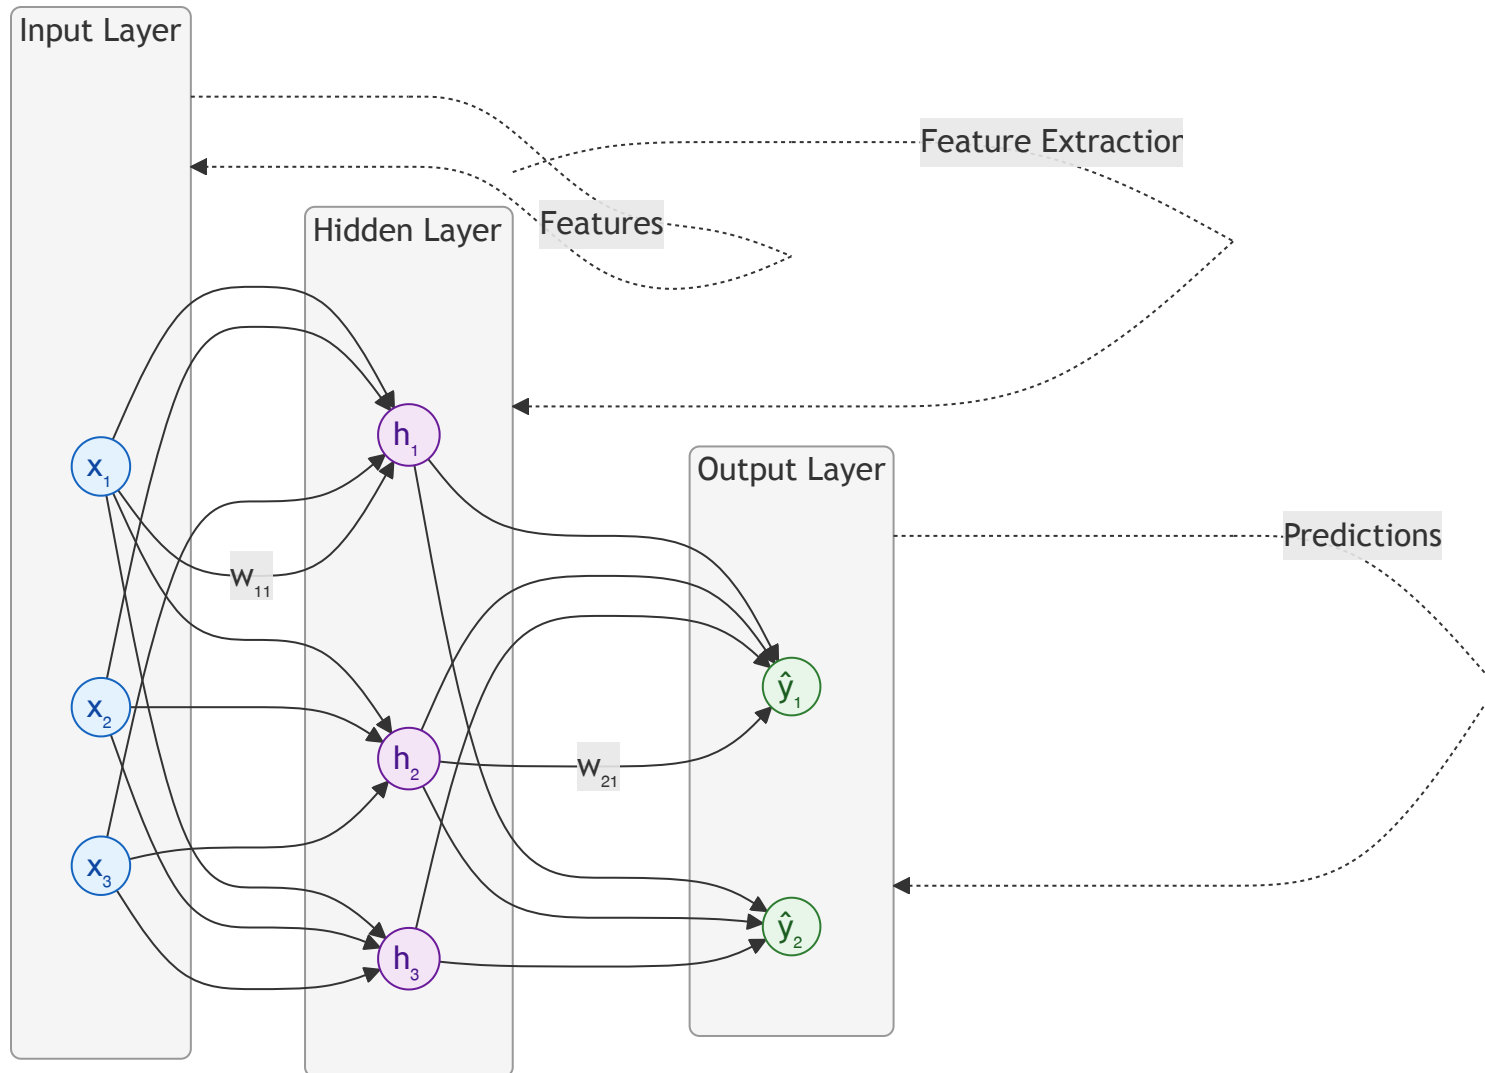



The Building Blocks of Deep Learning



The Building Blocks of Deep Learning

- Foundation of modern neural networks
- Versatile architecture for diverse problems
- Combines simplicity with powerful learning capabilities

From Neurons to Networks



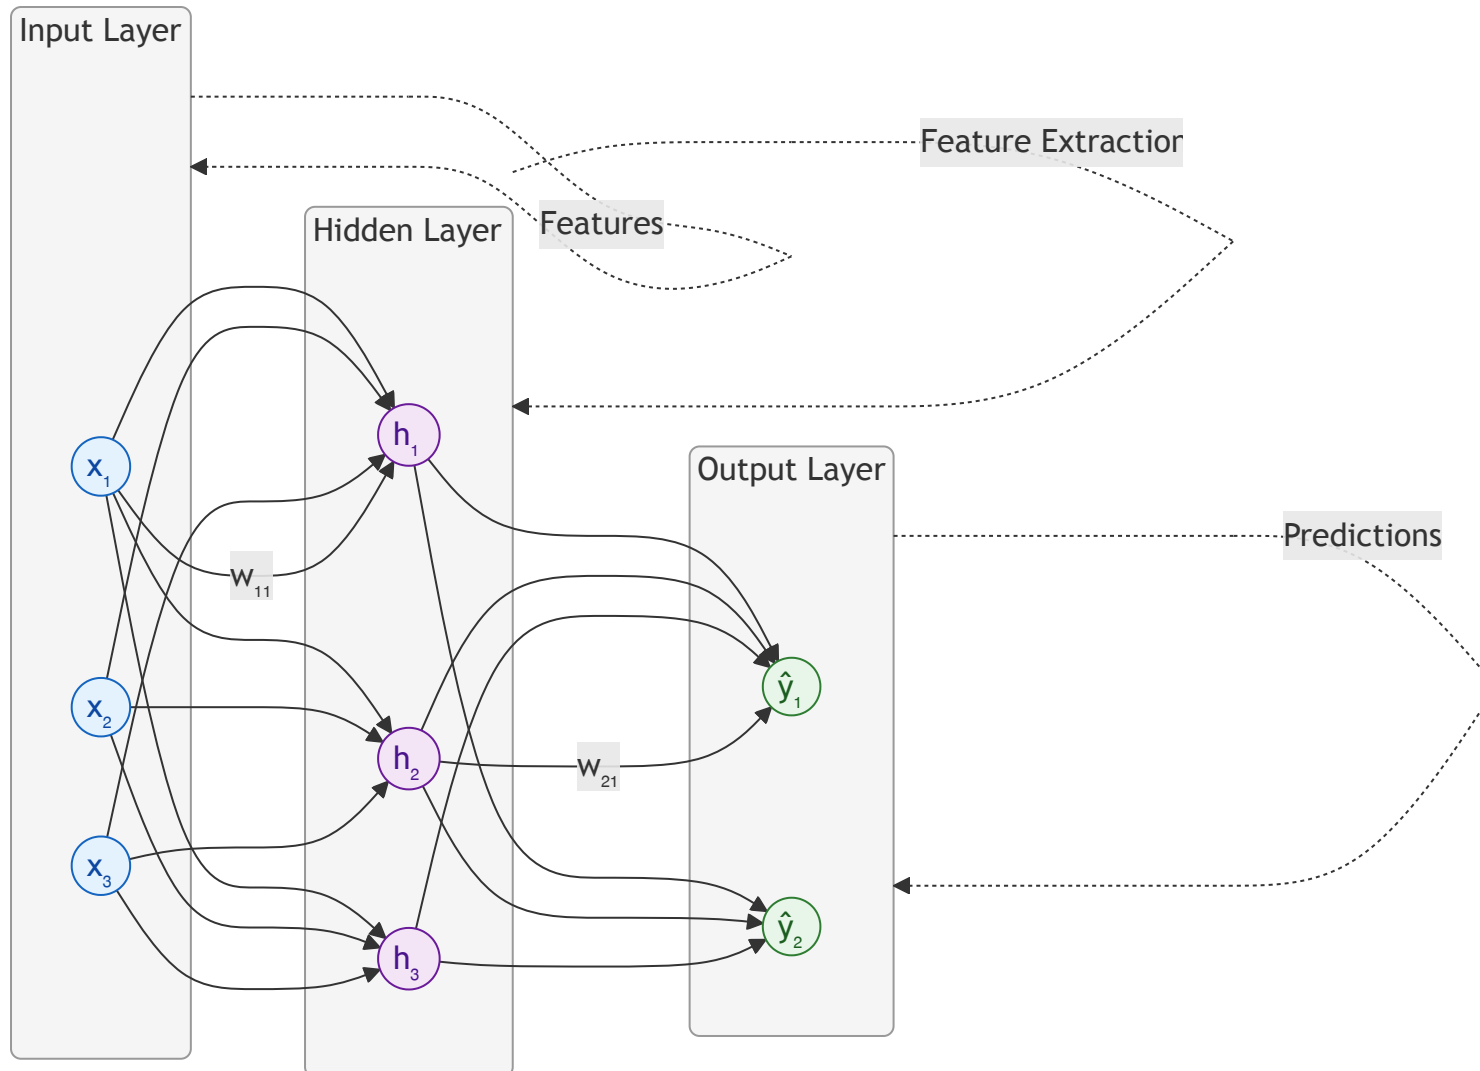
From Neurons to Networks

- **Biological inspiration:** Mimics brain's neural structure
 - Neurons receive, process, and transmit information
- **Artificial neuron:** Weighted sum + activation function
 - Processes inputs through mathematical operations
- **Network topology:** Input layer → Hidden layers → Output layer
 - Organized structure for information processing
- **Information flow:** Forward propagation for predictions
 - Data travels from input to output through the network

The Perceptron Journey

- **1958:** Rosenblatt's single-layer perceptron
 - First implementation of a neural learning algorithm
- **1969:** Minsky & Papert expose limitations (XOR problem)
 - Demonstrated that single-layer networks couldn't solve nonlinear problems
- **1986:** Rumelhart, Hinton & Williams introduce backpropagation
 - Breakthrough algorithm enabling training of multi-layer networks
- **Today:** Foundation for advanced architectures (CNNs, RNNs, Transformers)
 - Core concepts extended to specialized network designs


MLP Architecture



Key Components

- **Input layer:** Raw data reception
 - Receives and standardizes input features
- **Hidden layers:** Feature extraction and transformation
 - Learns hierarchical representations of data
- **Output layer:** Final prediction/classification
 - Produces the network's answer to the given problem

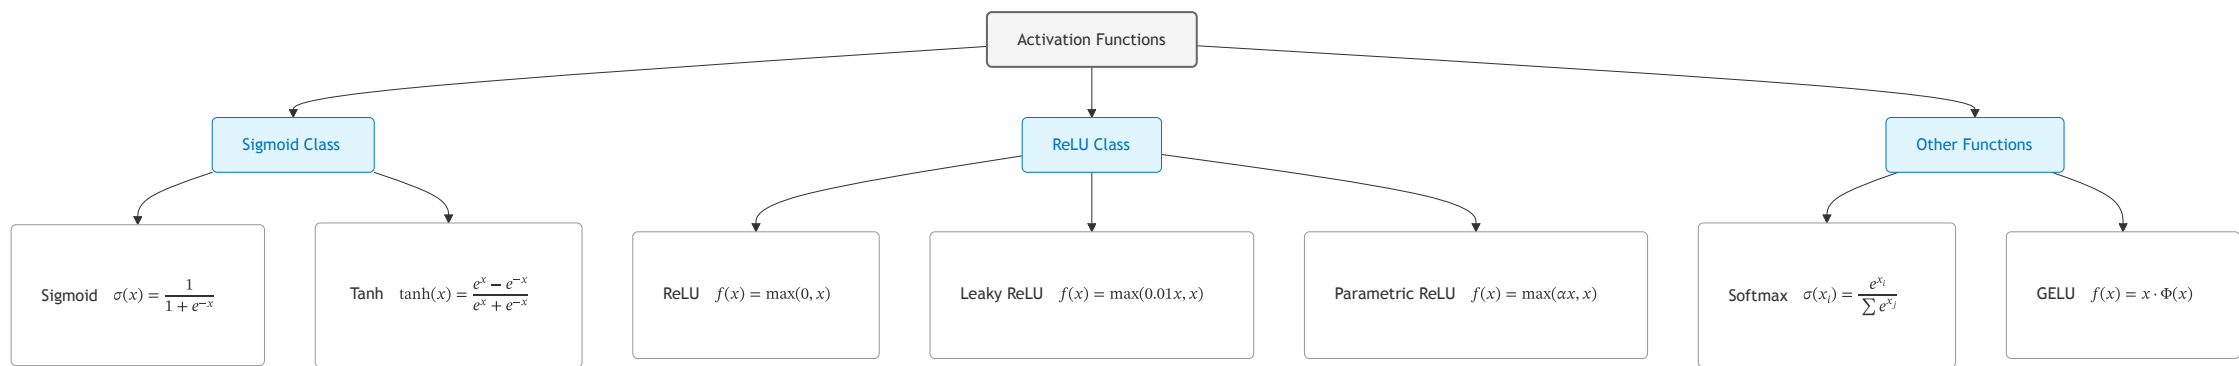
MLP Architecture



Key Components

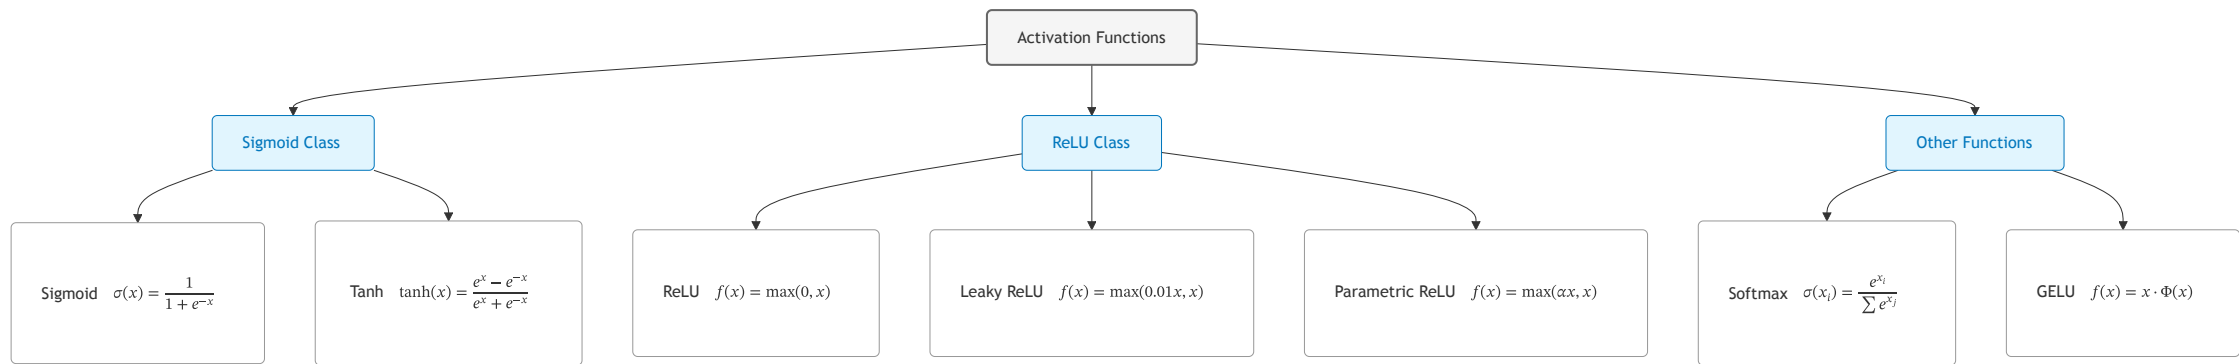
- **Weights & biases:** Learnable parameters
 - Adjusted during training to minimize error
- **Activation functions:** Introduce non-linearity
 - Enable the network to learn complex patterns

Activation Functions



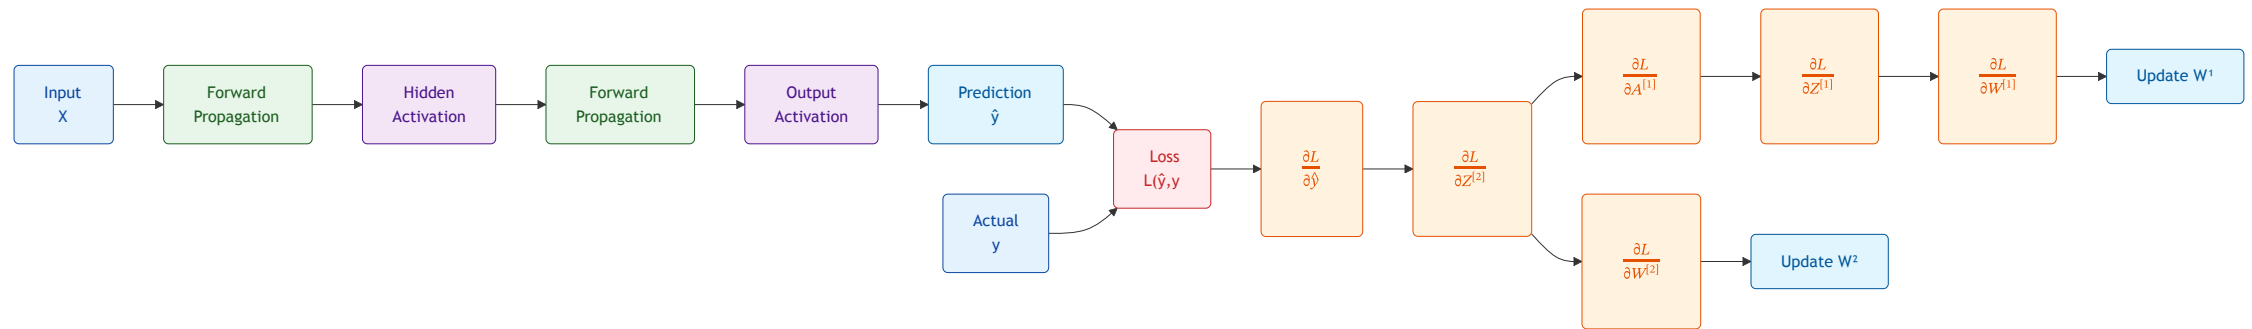
Function	Formula	Characteristics
Sigmoid	$\sigma(x) = \frac{1}{1+e^{-x}}$	Output range [0,1], vanishing gradient
Tanh	$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$	Output range [-1,1], zero-centered

Activation Functions



Function	Formula	Characteristics
ReLU	$f(x) = \max(0, x)$	Computationally efficient, sparse activation
Leaky ReLU	$f(x) = \max(0.01x, x)$	Prevents dying ReLU problem

Forward Propagation



For each layer l :

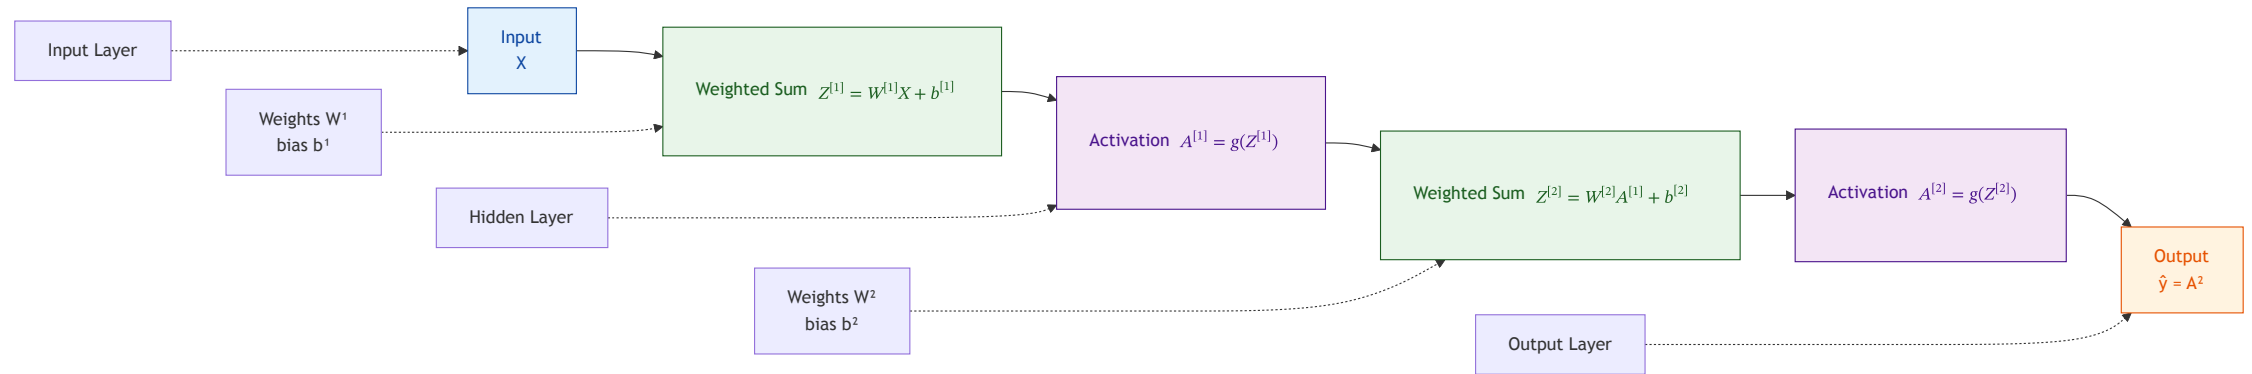
$$Z^{[l]} = W^{[l]} \cdot A^{[l-1]} + b^{[l]}$$

$$A^{[l]} = g^{[l]}(Z^{[l]})$$

Where:

- $W^{[l]}$ = weights matrix
- $b^{[l]}$ = bias vector


Forward Propagation



Where:

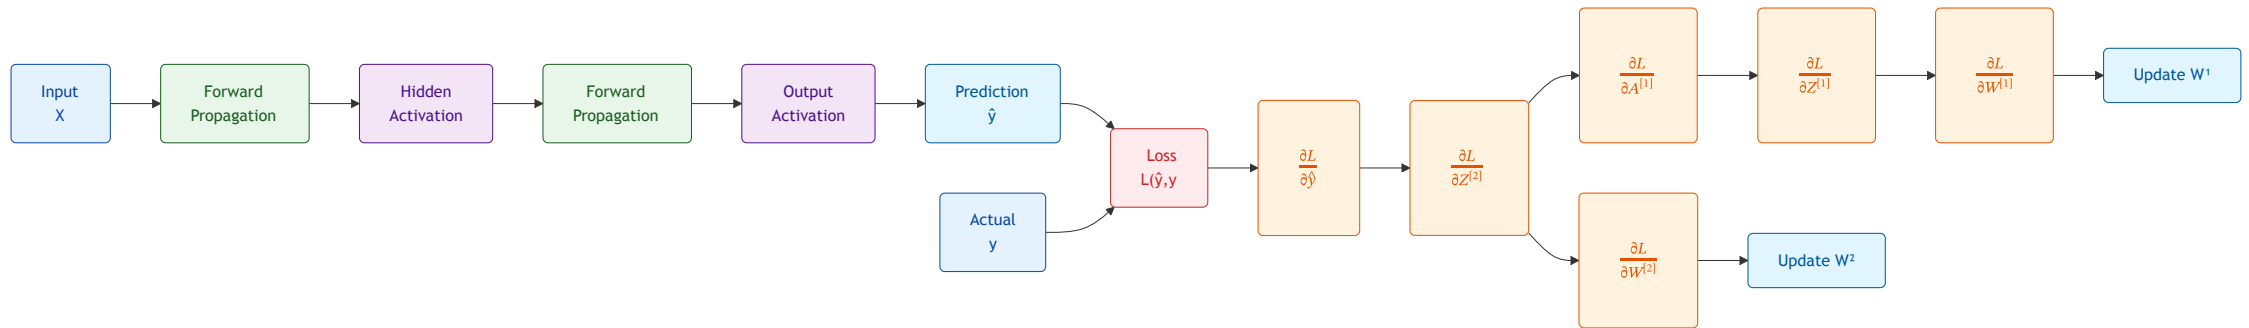
- $g^{[l]}$ = activation function
- $A^{[l]}$ = activation output

Backpropagation: Learning Process



1. **Forward pass:** Compute predictions
 - Process inputs through the network
2. **Error calculation:** Compare with ground truth

Backpropagation: Learning Process



3. Backward pass: Compute gradients

4. Parameter update: Adjust weights and biases

$$W^{[l]} = W^{[l]} - \alpha \frac{\partial J}{\partial W^{[l]}}$$

$$b^{[l]} = b^{[l]} - \alpha \frac{\partial J}{\partial b^{[l]}}$$

Loss Functions

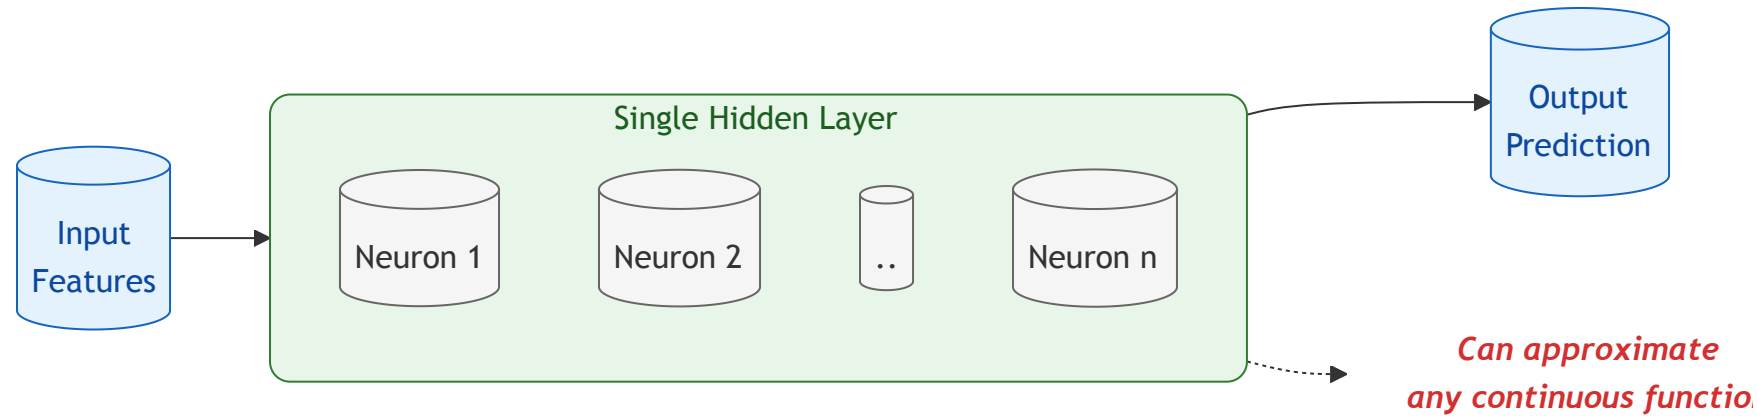
Task	Loss Function	Formula
Regression	Mean Squared Error	$\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$
Binary Classification	Binary Cross-Entropy	$-\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$

Loss Functions

Task	Loss Function	Formula
Multi-class Classification	Categorical Cross-Entropy	$-\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^m y_{ij} \log(\hat{y}_{ij})$

- Loss guides the learning process
- Different tasks use specialized error measurements
- Optimization aims to minimize loss

Universal Approximation Theorem

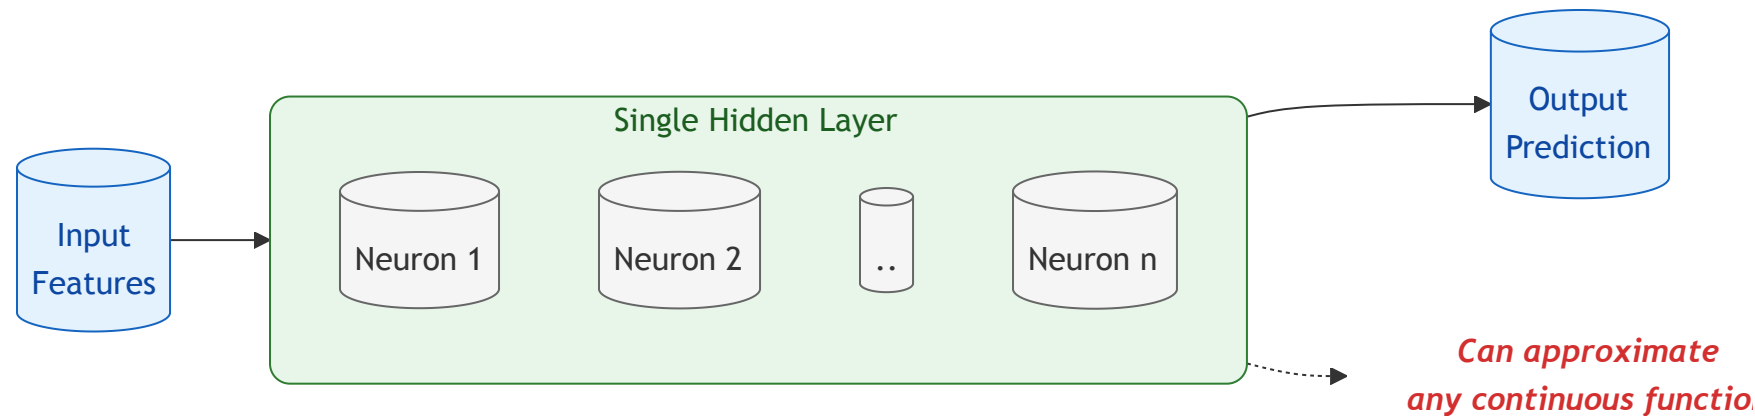


Universal Approximation Theorem

"A feedforward network with a single hidden layer containing a finite number of neurons can approximate any continuous function, under mild assumptions on the activation function."

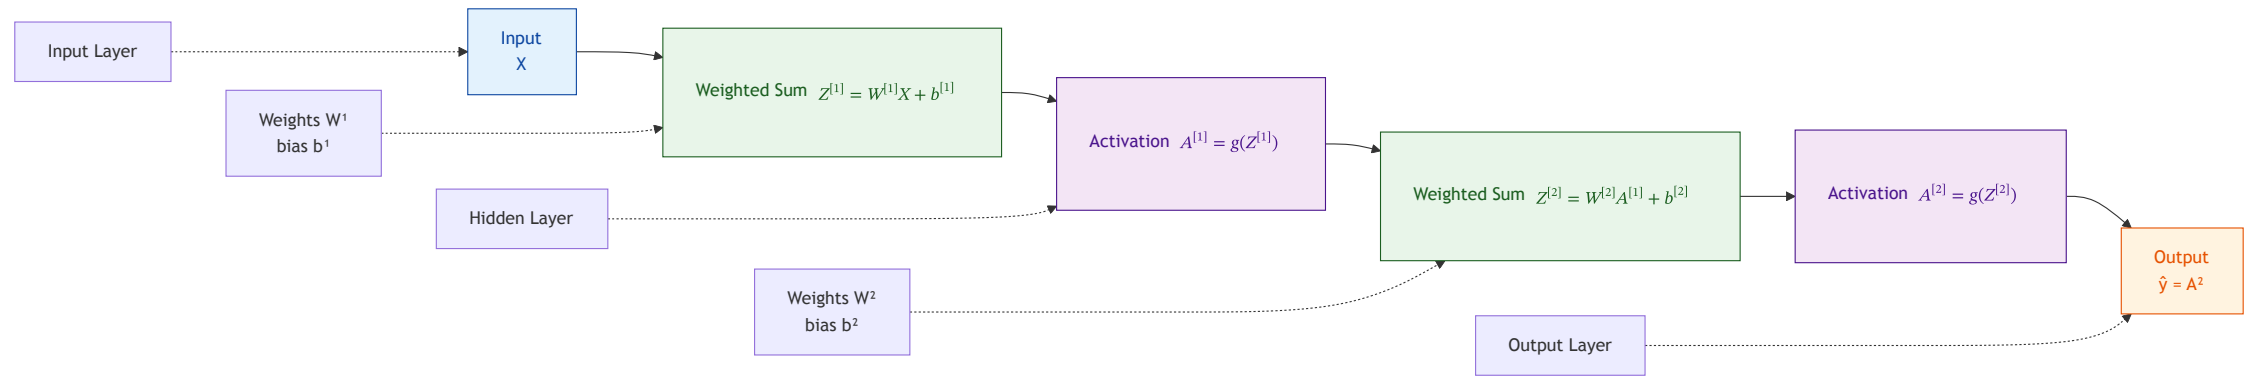
- The theoretical foundation for MLP capabilities

Universal Approximation Theorem



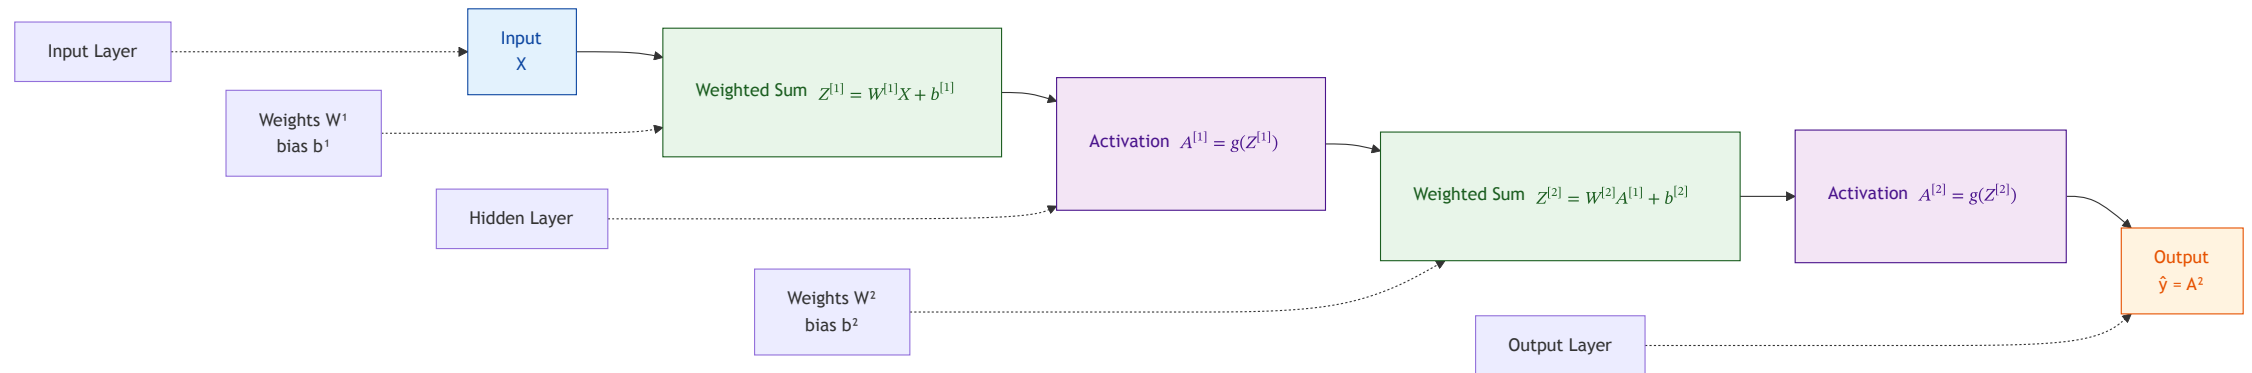
- More complex functions may require more neurons
- Practical implementations must balance capacity and training challenges

Visualizing Decision Boundaries



- **Linear boundaries:** Single-layer perceptrons
 - Separate data with straight lines
- **Non-linear boundaries:** MLPs with hidden layers
 - Can form complex separation surfaces

Visualizing Decision Boundaries



- **Complexity increases** with deeper architectures
- Explore an interactive demo at perceptron.marcr.xyz

Quick Quiz: Test Your Knowledge

Which of these problems can a single-layer perceptron solve?

- A) XOR problem
- B) Linear classification
- C) Image recognition
- D) All of the above

Use the poll feature to submit your answer!

Practical Implementation Challenges

What's your biggest challenge with neural networks?

- Understanding the math
- Choosing the right architecture
- Overfitting/underfitting
- Computational resources
- Interpreting results

Share your thoughts!

Thank You

Contact Information

- Email: hi@marcr.xyz

Resources

- Interactive Demo: perceptron.marcr.xyz
- Slides: github.com/mabreyes/dlsu-lecture-slides